

Programmable Networks with Synthesis



Ahmed ElHassany



Petar Tsankov



Laurent Vanbever



Martin Vechev

Network Misconfigurations are Common

Amazon server outage affects millions of companies and causes online chaos

MARCH 1, 2017 12:39PM

AMAZON'S giant servers crashed today causing chaos for millions of companies and people that use the cloud service and affected everything larger web sites to people's smarthomes and even library catalogues.

Amazon's massive AWS outage was caused by human error

One incorrect command and the whole internet suffers.

BY JASON DEL REY | @DELREY | MAR 2, 2017, 2:20PM EST

[TWEET](#) [SHARE](#) [LINKEDIN](#)

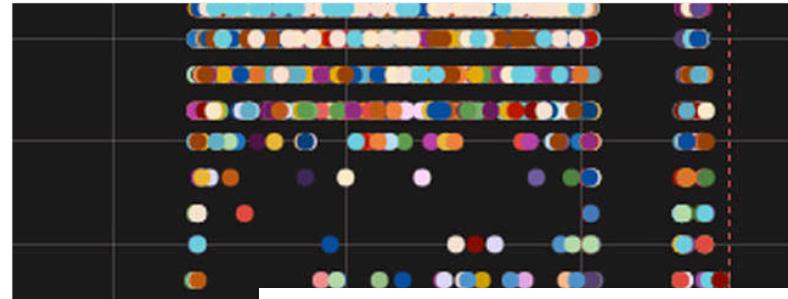


Sean Gallup / Getty

Amazon today blamed human error for the big AWS outage that took down a bunch of large Internet sites for several hours on Tuesday afternoon.

CloudFlare apologizes for Telia screwing you over

Unhappy about massive outage

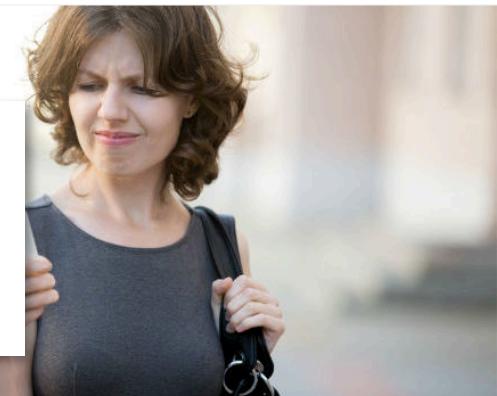


Networks

Level3 switch config blunder blamed for US-wide VoIP blackout

Network dropped calls because it was told to

21 Jun 2016 at 20:34, Kieren Mc



The summer of network misconfigurations

by Joanne Godfrey on August 11, 2016 in Application Connectivity Management, Firewall Change Management, Information Security, Risk Management and Vulnerabilities, Security Policy Management

[EMAIL](#)

[Share](#)

32

[Tweet](#)

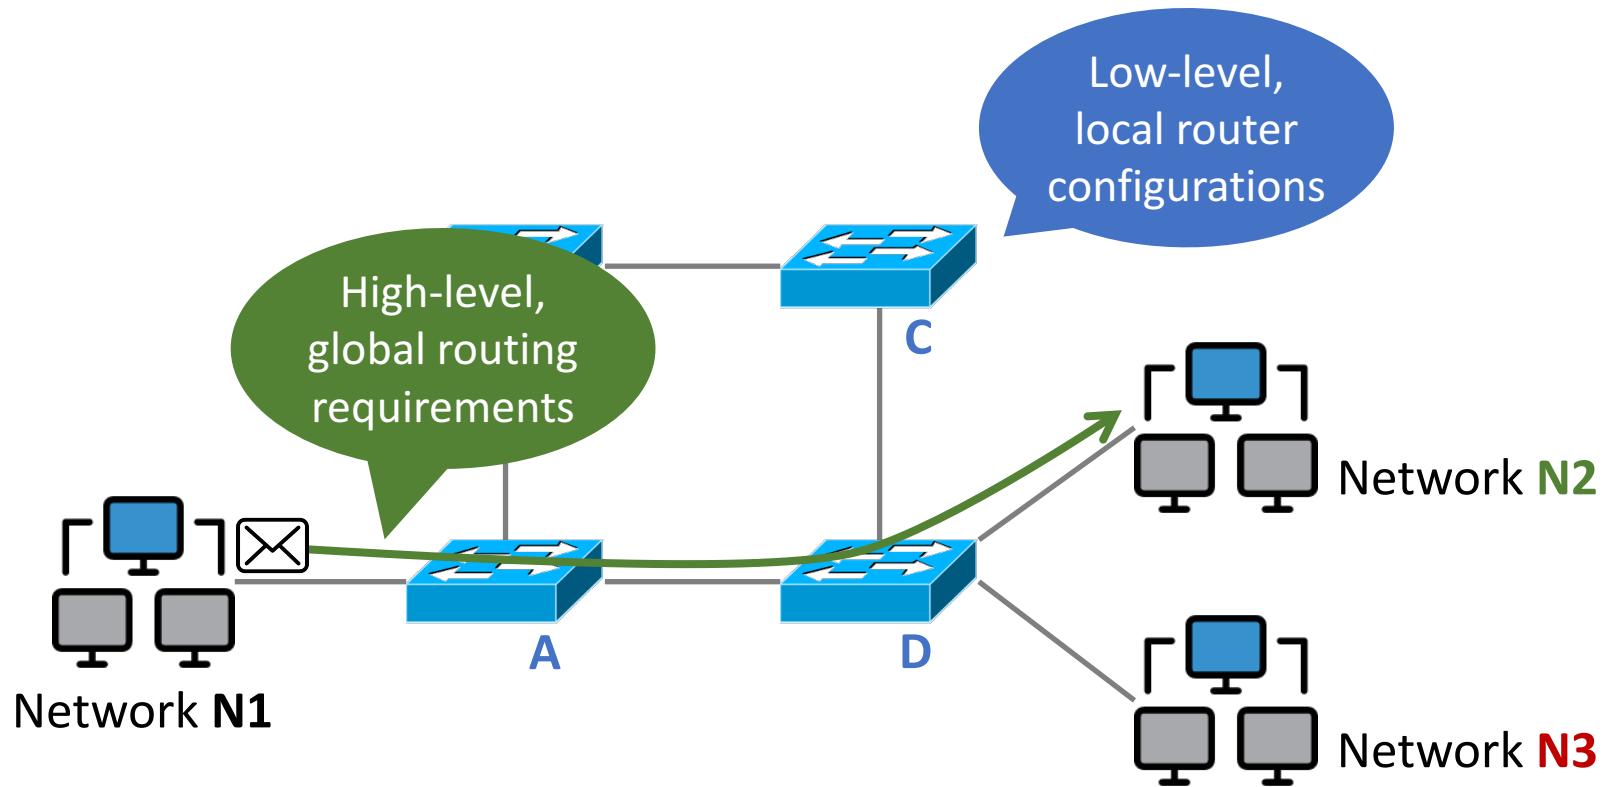
[Share 6](#)

[Like 6](#)

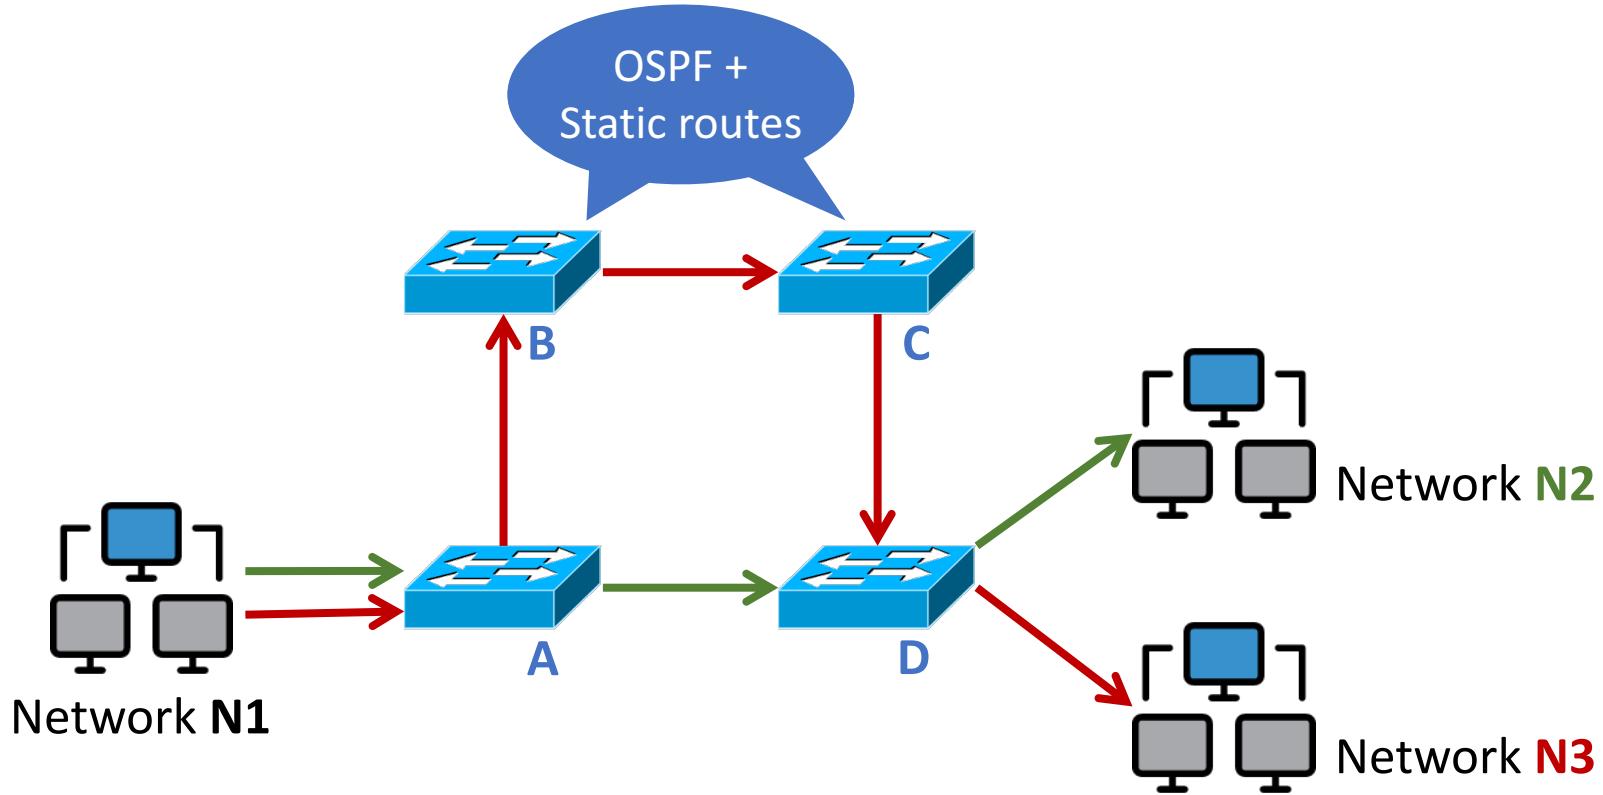
5 Oct 2016 at 11:33, Shaun Nichols

[Reddit](#) [Twitter](#) [Facebook](#) [LinkedIn](#)

What Makes Network Configuration Hard?



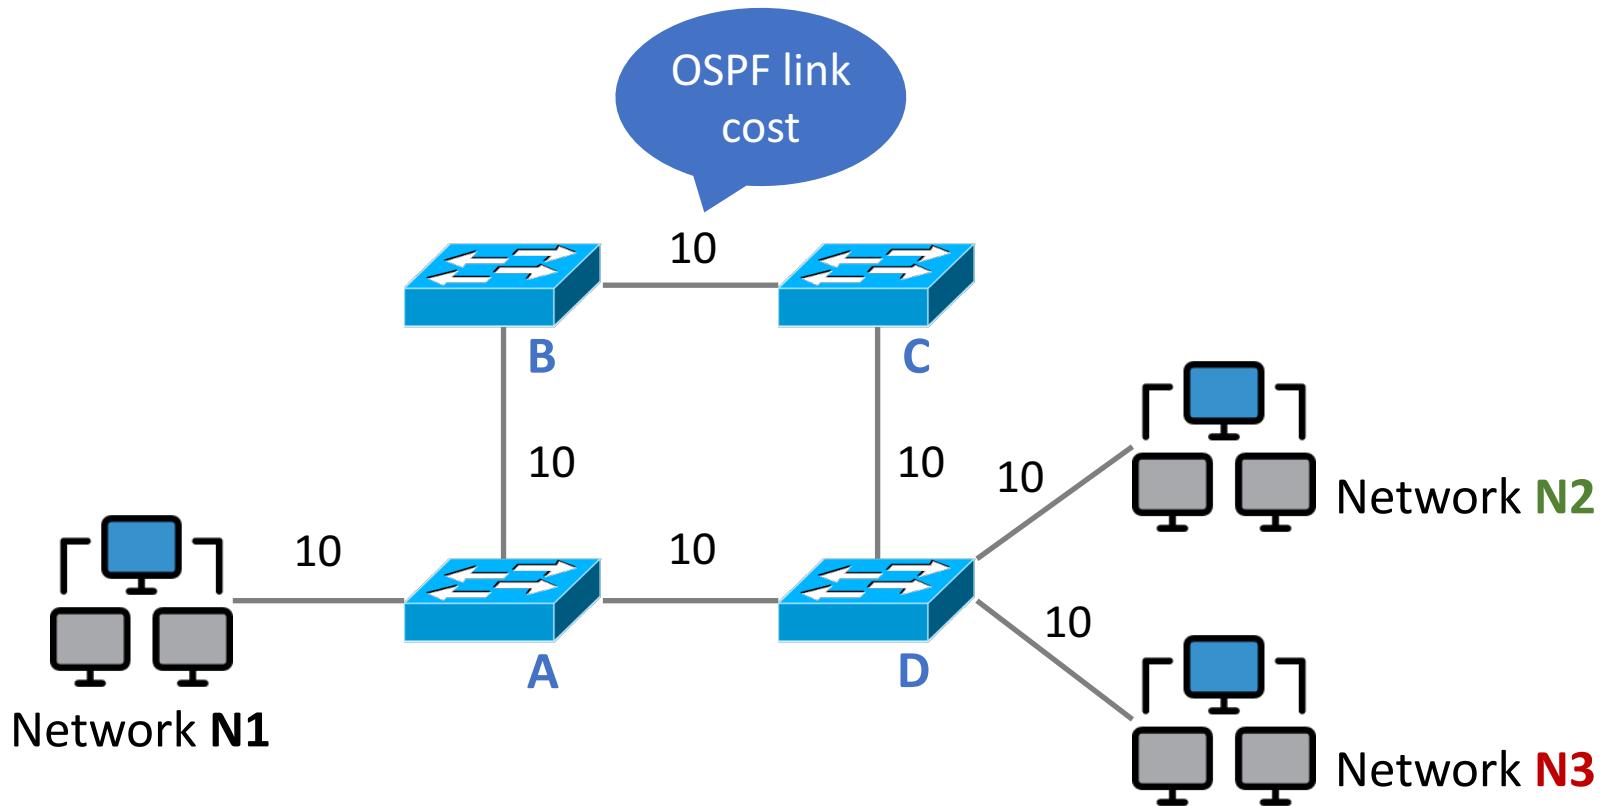
Example



R1: Packets from **N1** to **N2** must follow the path **A** → **D**

R2: Packets from **N1** to **N3** must follow the path **A** → **B** → **C** → **D**

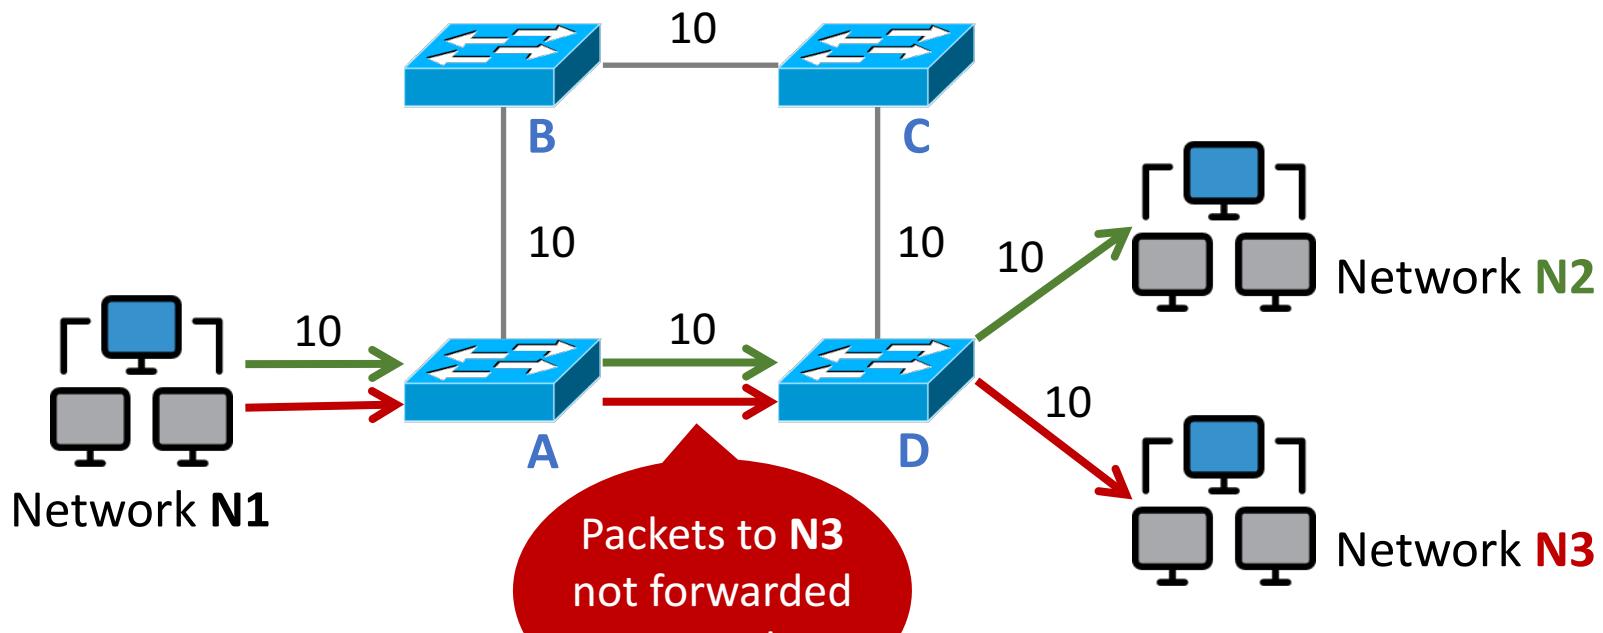
Example



R1: Packets from **N1** to **N2** must follow the path **A → D**

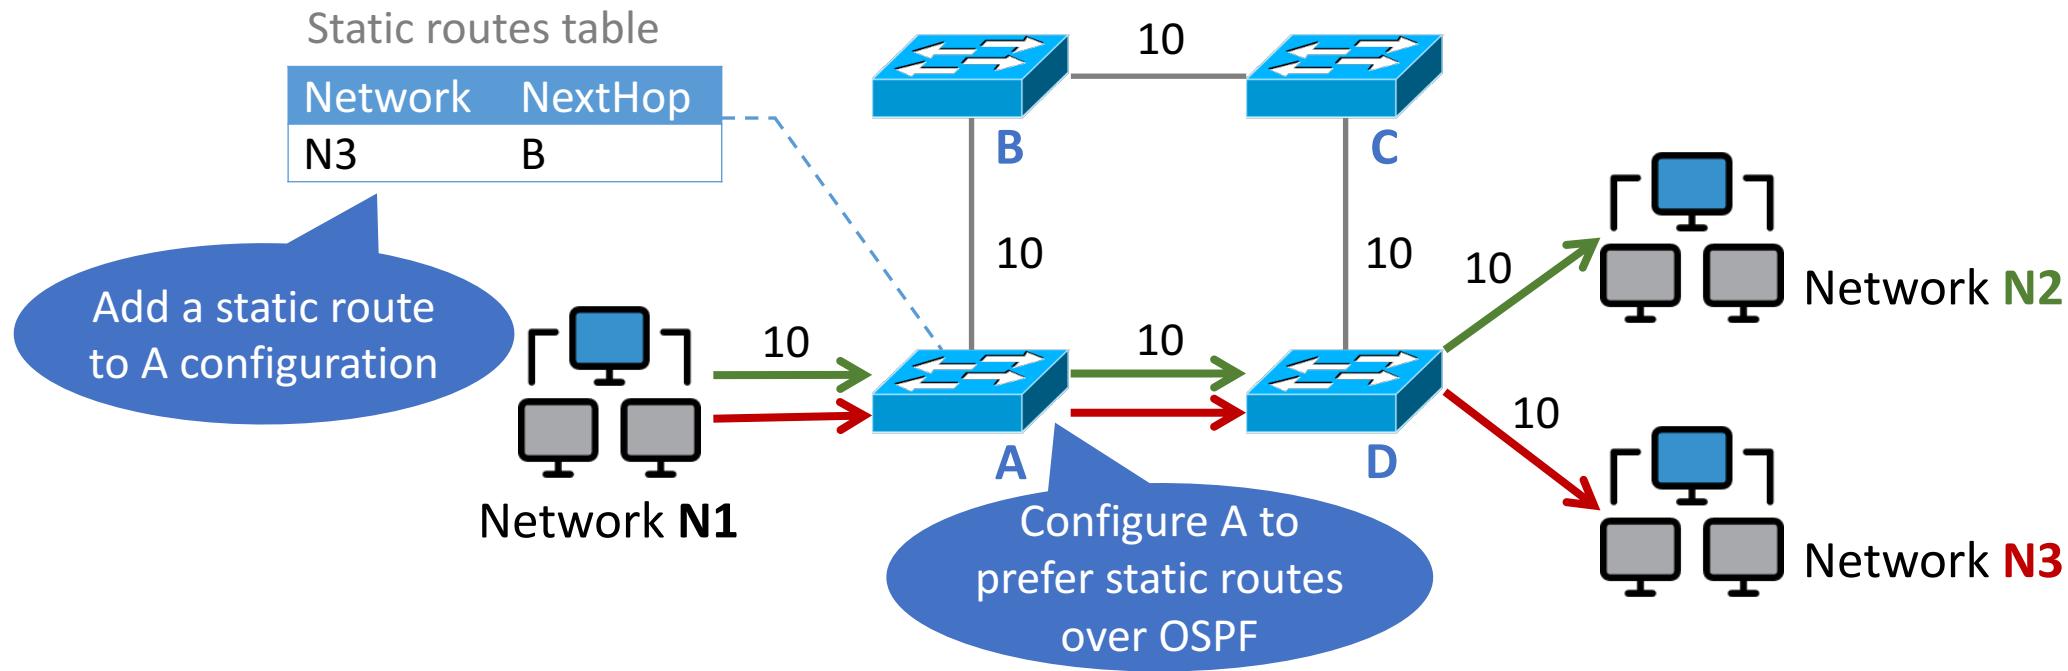
R2: Packets from **N1** to **N3** must follow the path **A → B → C → D**

Example



- ✓ R1: Packets from **N1** to **N2** must follow the path **A → D**
- ✗ R2: Packets from **N1** to **N3** must follow the path **A → B → C → D**

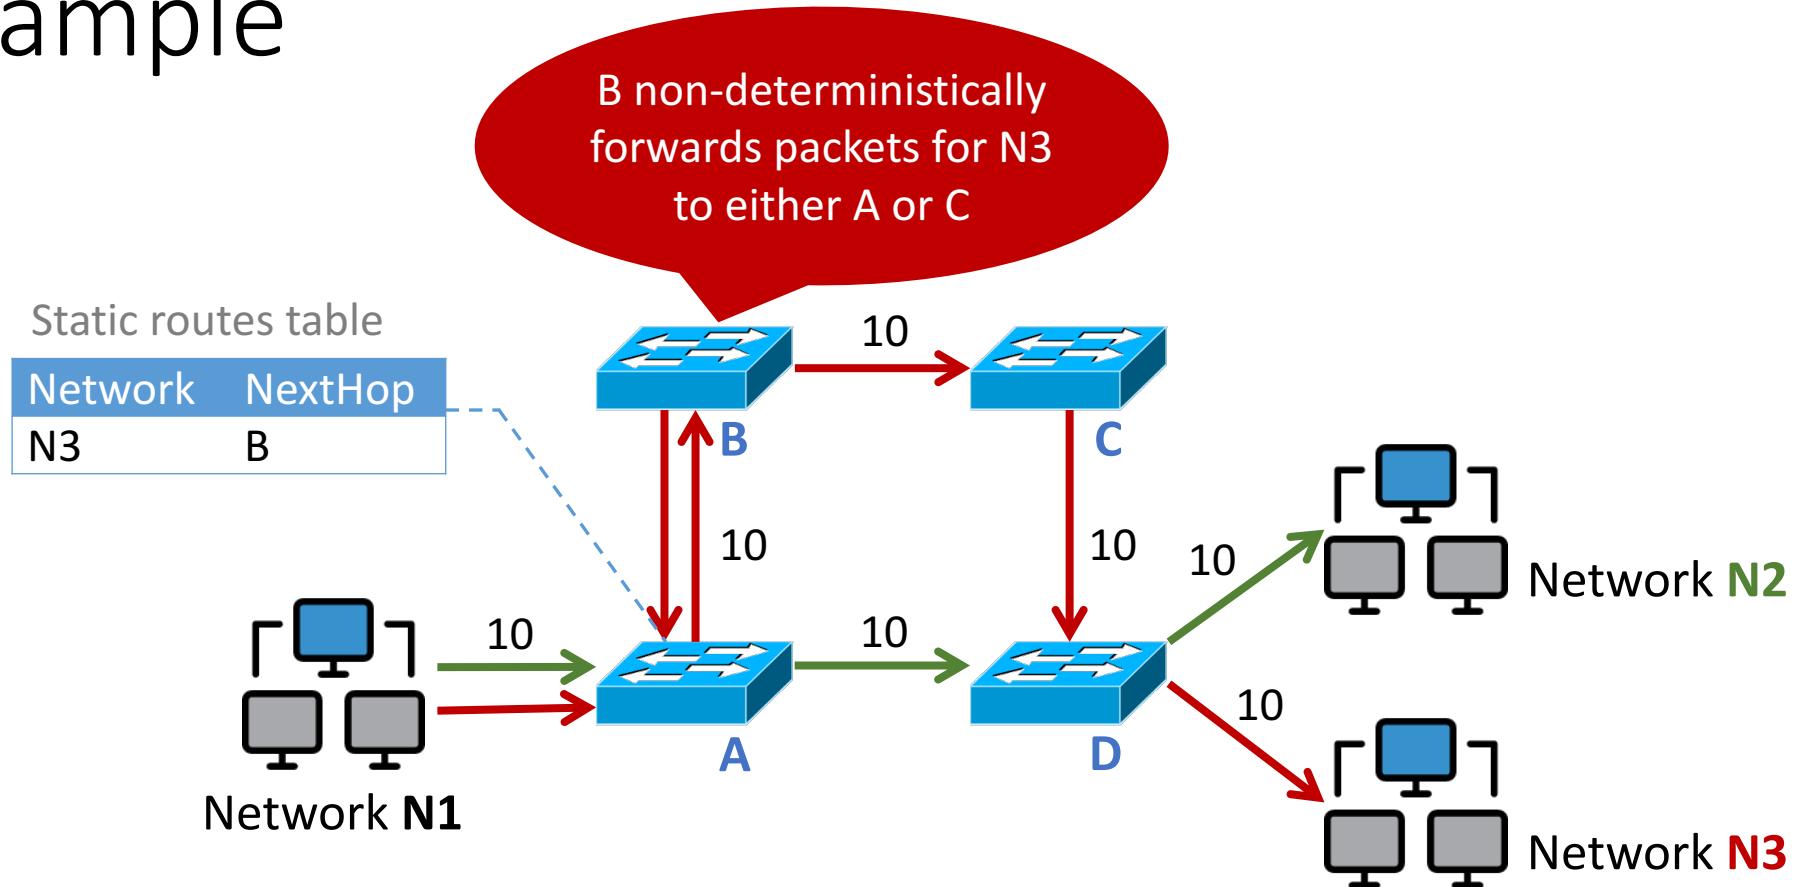
Example



✓ R1: Packets from **N1** to **N2** must follow the path **A → D**

✗ R2: Packets from **N1** to **N3** must follow the path **A → B → C → D**

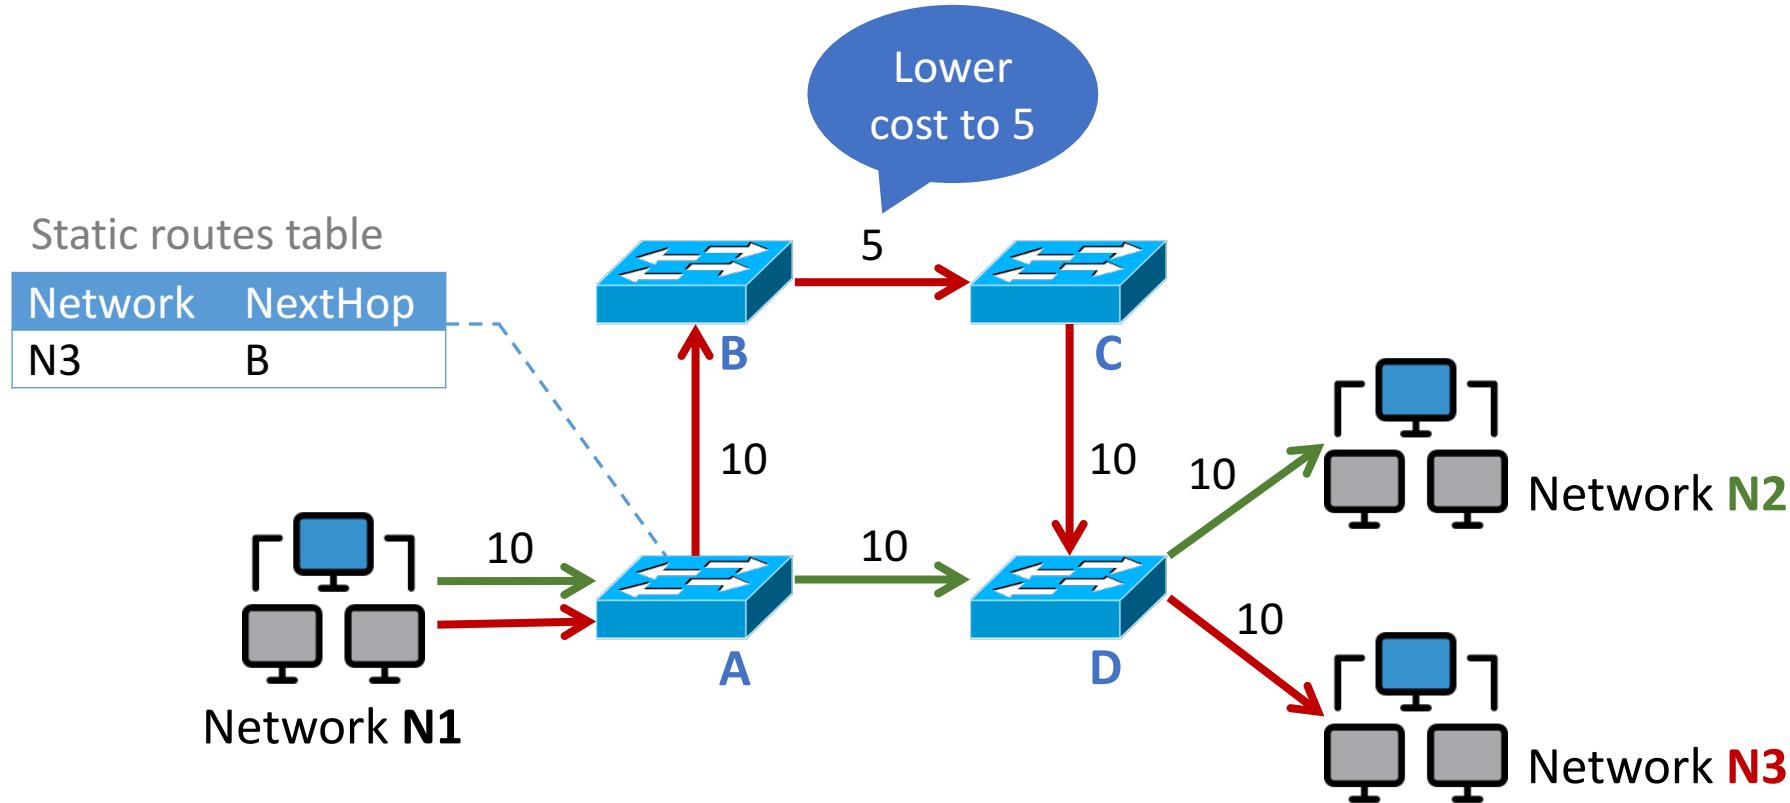
Example



✓ R1: Packets from **N1** to **N2** must follow the path **A** → **D**

✗ R2: Packets from **N1** to **N3** must follow the path **A** → **B** → **C** → **D**

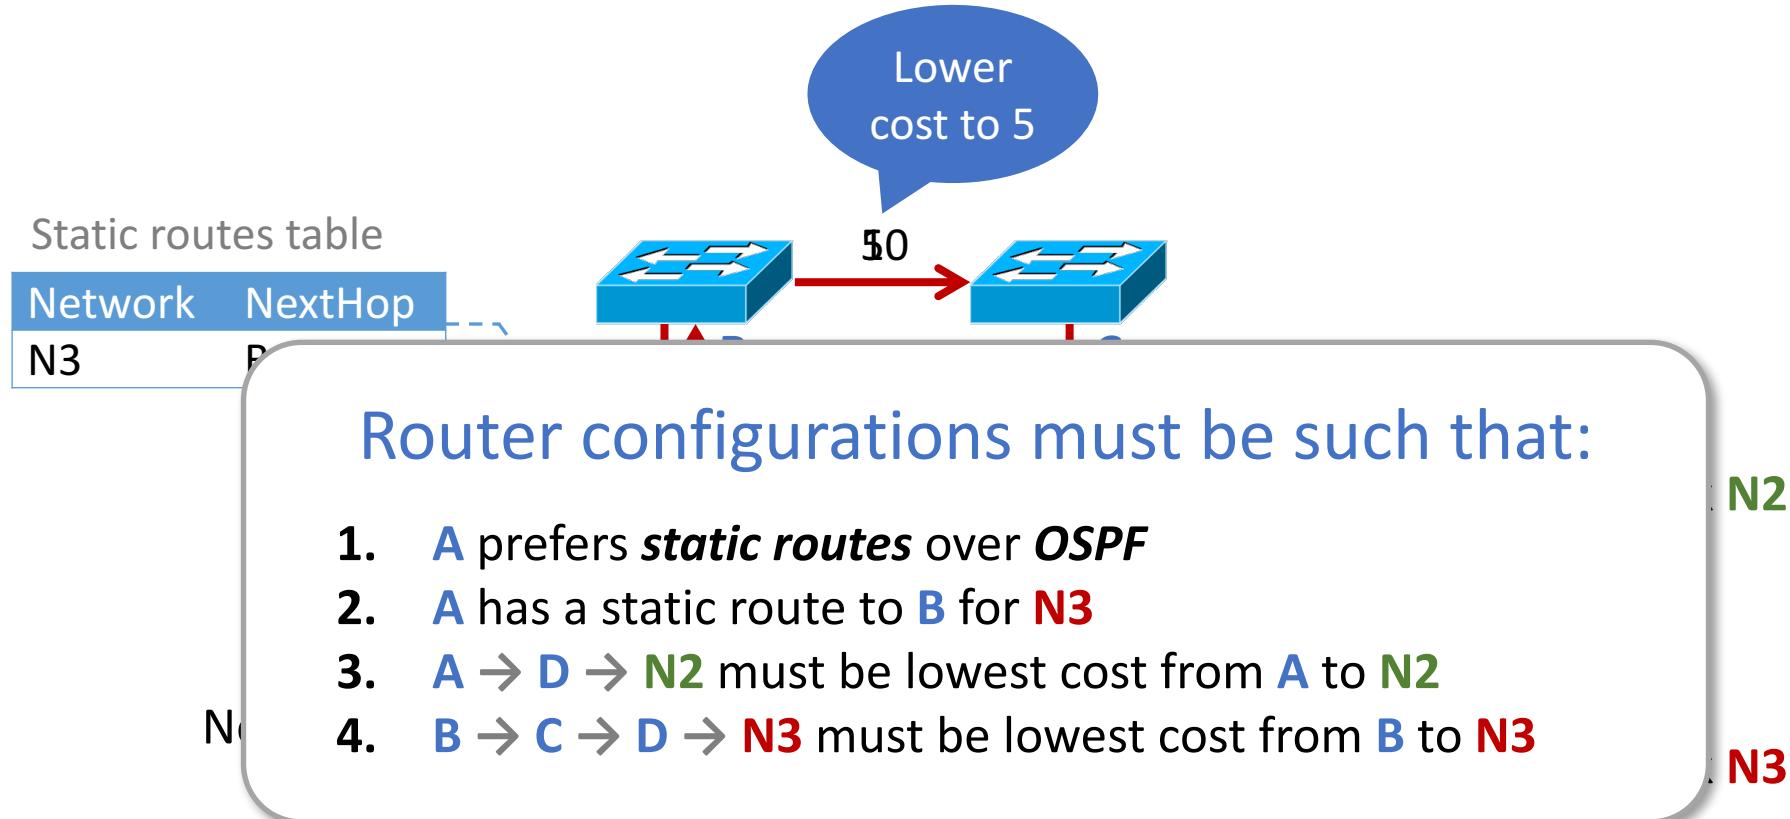
Example



✓ R1: Packets from **N1** to **N2** must follow the path **A → D**

✗ R2: Packets from **N1** to **N3** must follow the path **A → B → C → D**

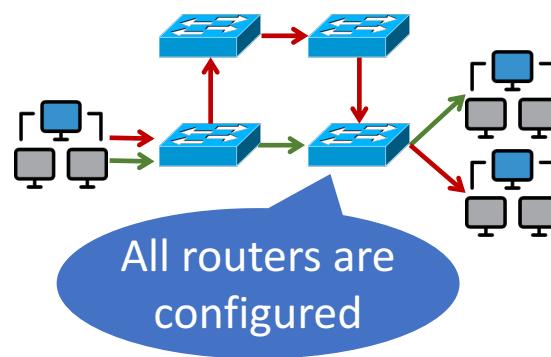
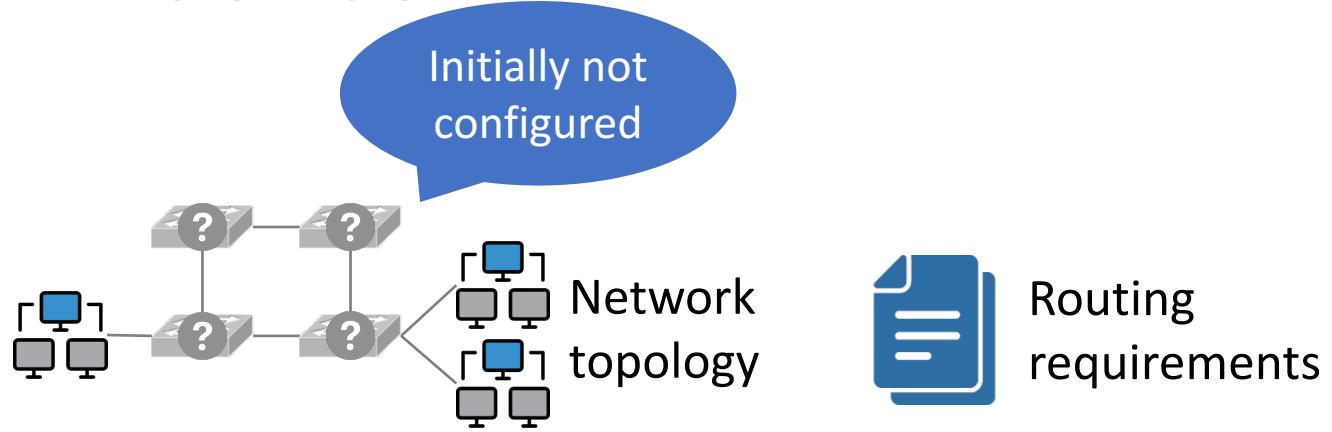
Example



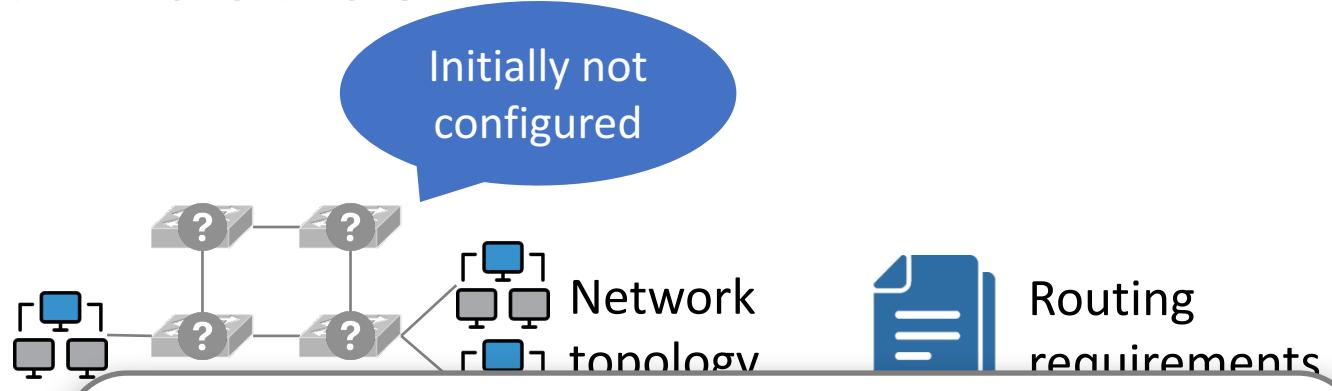
✓ R1: Packets from **N1** to **N2** must follow the path **A → D**

✗ R2: Packets from **N1** to **N3** must follow the path **A → B → C → D**

Current Practice

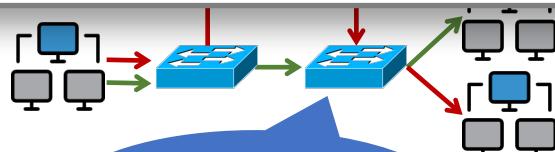


Current Practice



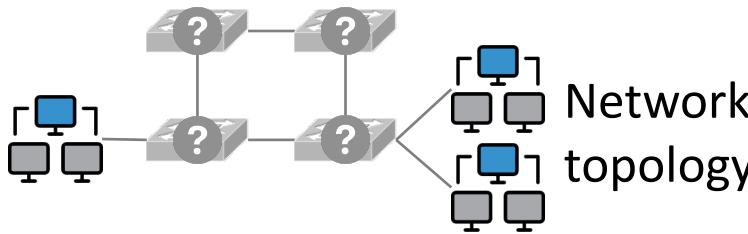
Problems and Challenges

- *Diversity* in protocol expressiveness
- Protocol *dependencies*
- No *correctness* guarantees



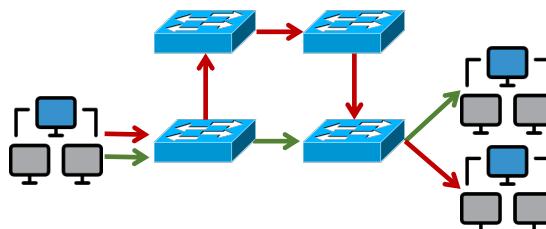
All routers are configured

Wanted: Programmable Networks with Synthesis

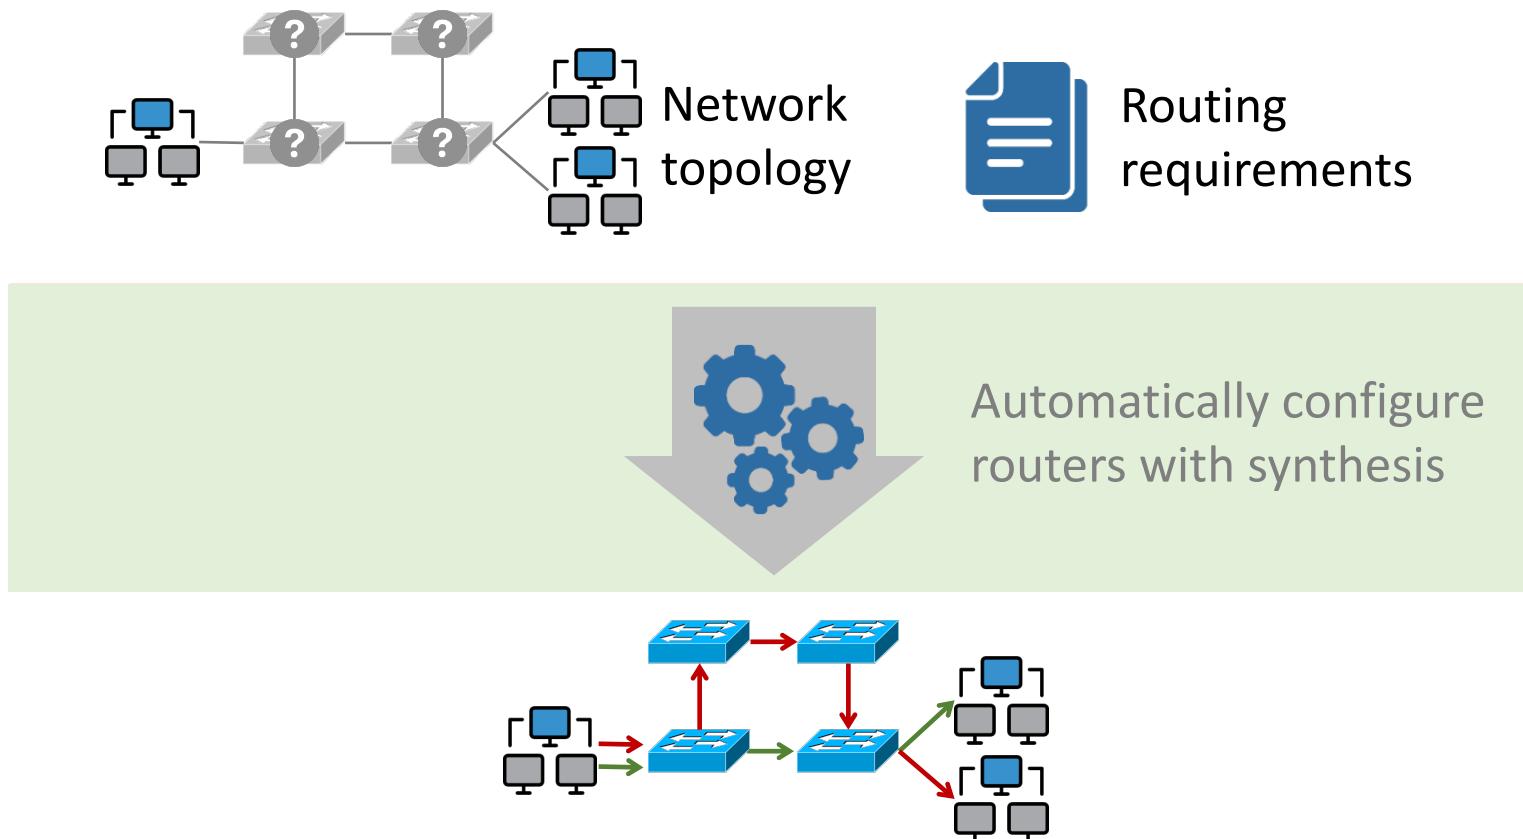


Routing
requirements

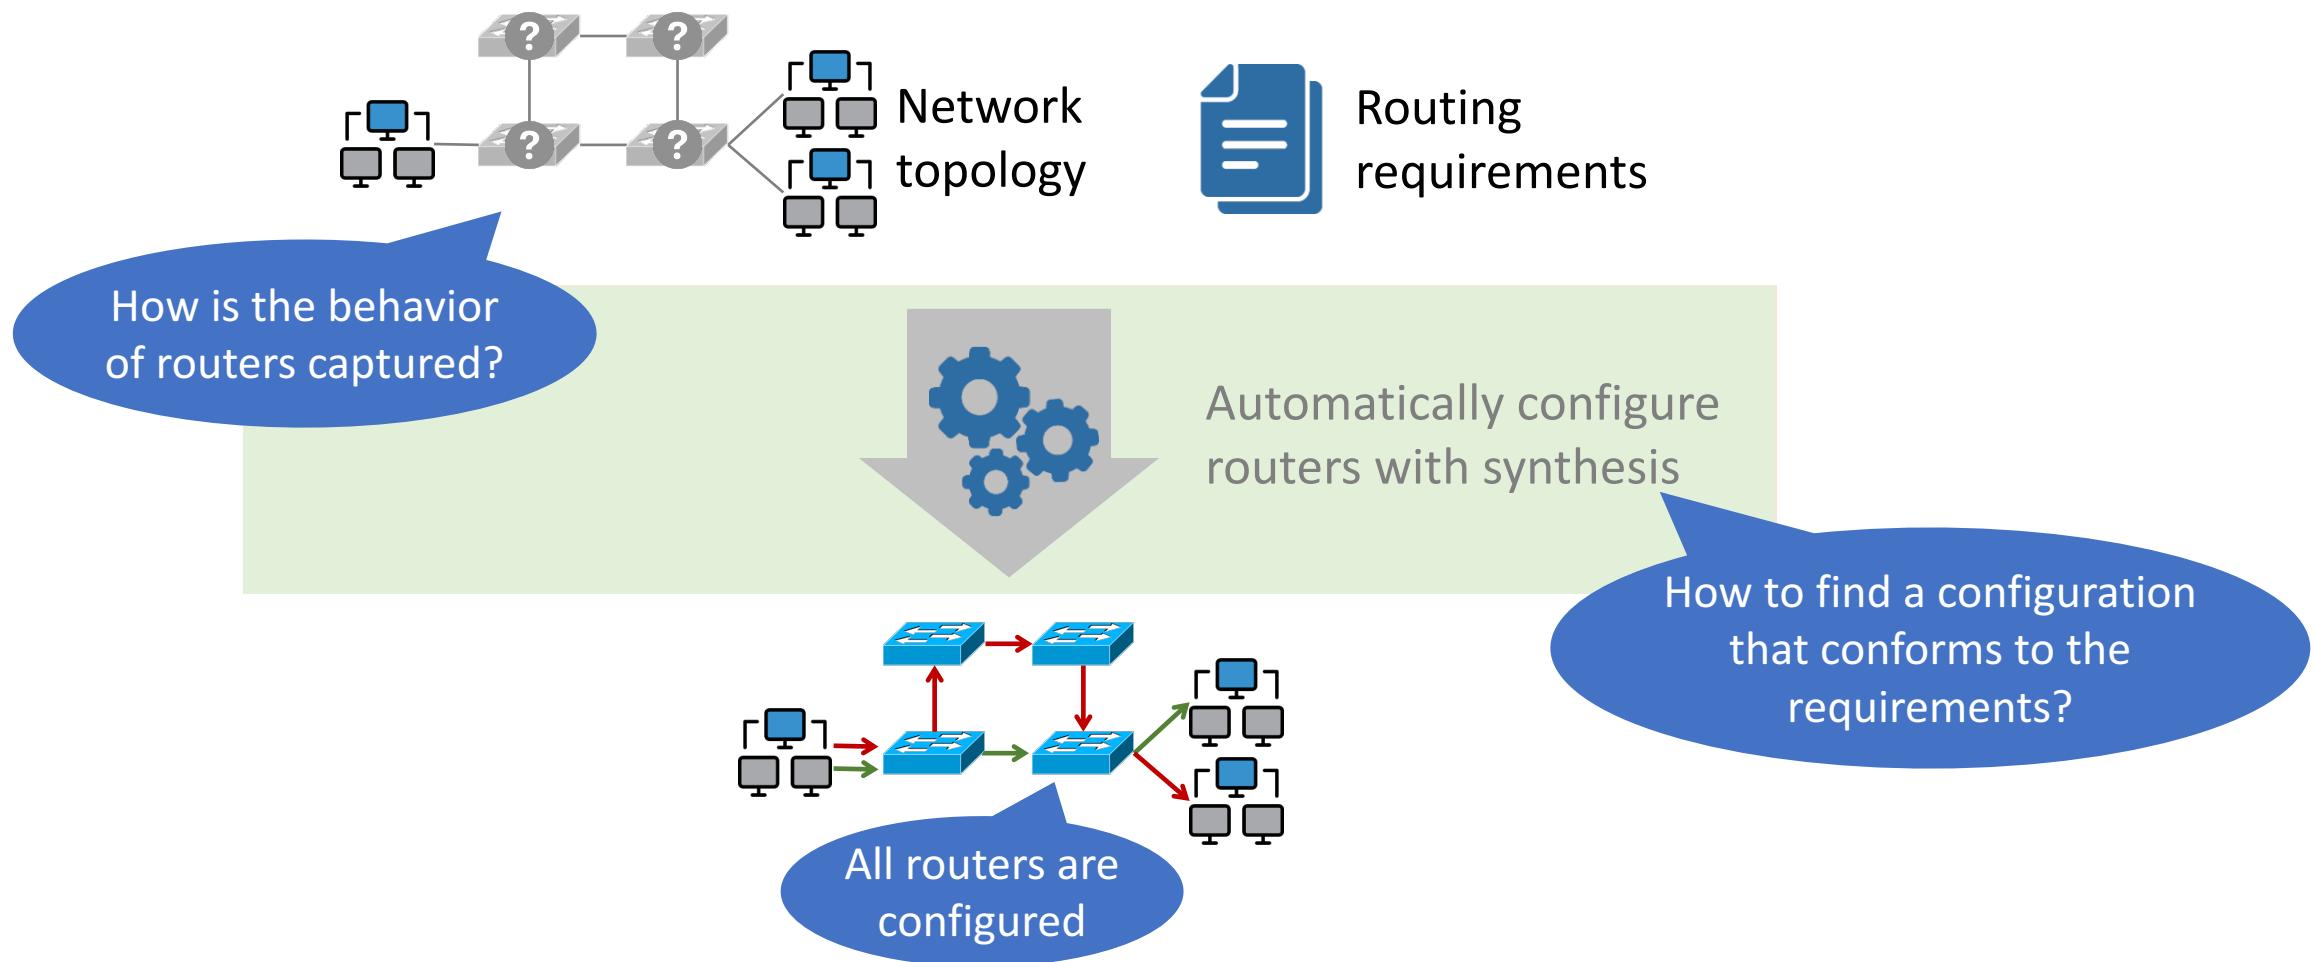
Operators manually
configure each router



Wanted: Programmable Networks with Synthesis



Wanted: Programmable Networks with Synthesis



Programmable Networks with *Synthesis*: Dimensions

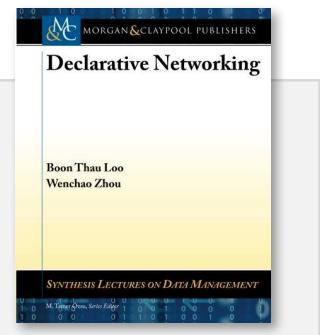
Deployment scenarios	Datacenter Incremental	ISP Enterprise
Routing protocols	Deterministic OSPF BGP Static routes MPLS	Probabilistic ECMP Gossip
Requirements	Paths Isolation Reachability Waypointing	Failures Congestion
Synthesis Techniques	Enumerative learning Symbolic execution (SyNET*)	Probabilistic Constraint solving CEGIS

***SyNET**: <http://synet.ethz.ch>

Capturing Network Behavior



Key idea: Express routing protocols, along with their dependencies, in *stratified Datalog*



Datalog (Graph Reachability)

Input

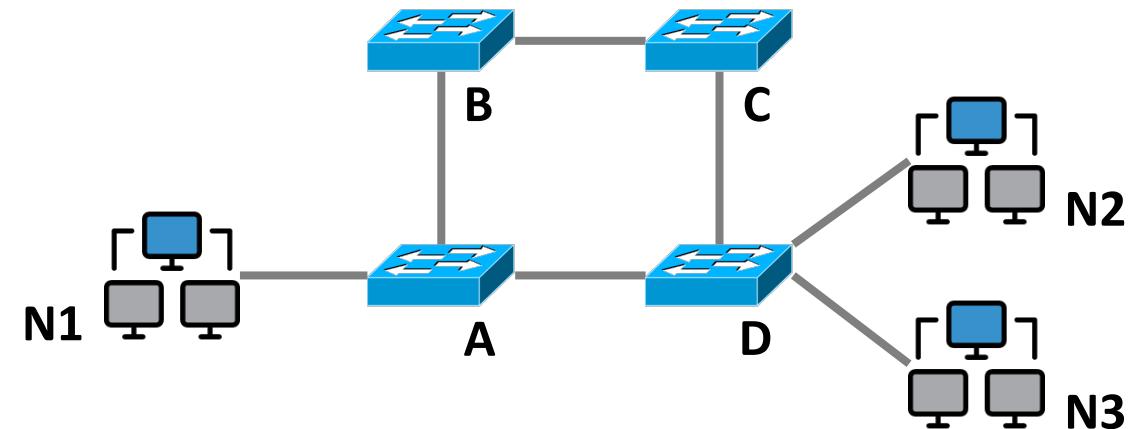
$link(n1, a)$
 $link(a, b)$
...

Program

$path(X, Y) \leftarrow link(X, Y)$
 $path(X, Y) \leftarrow link(X, Z), path(Z, Y)$

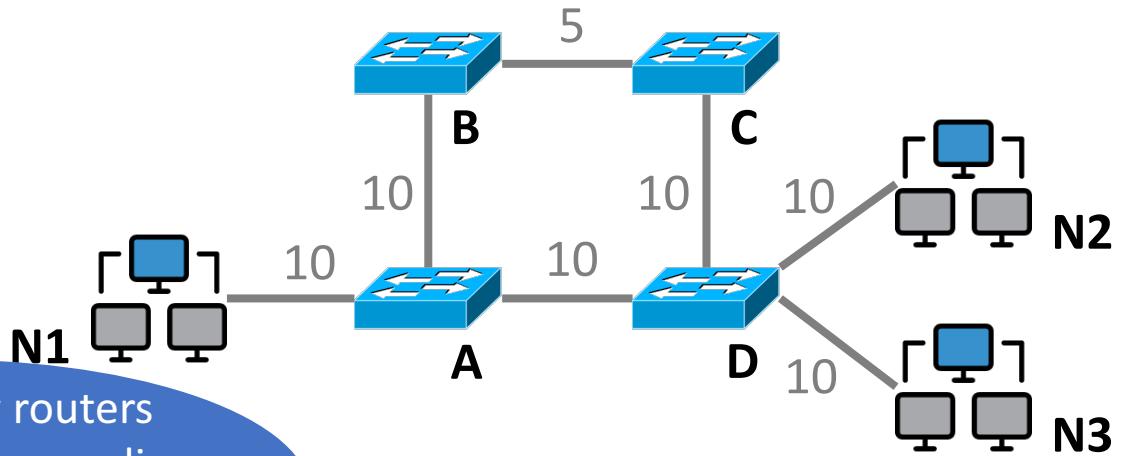
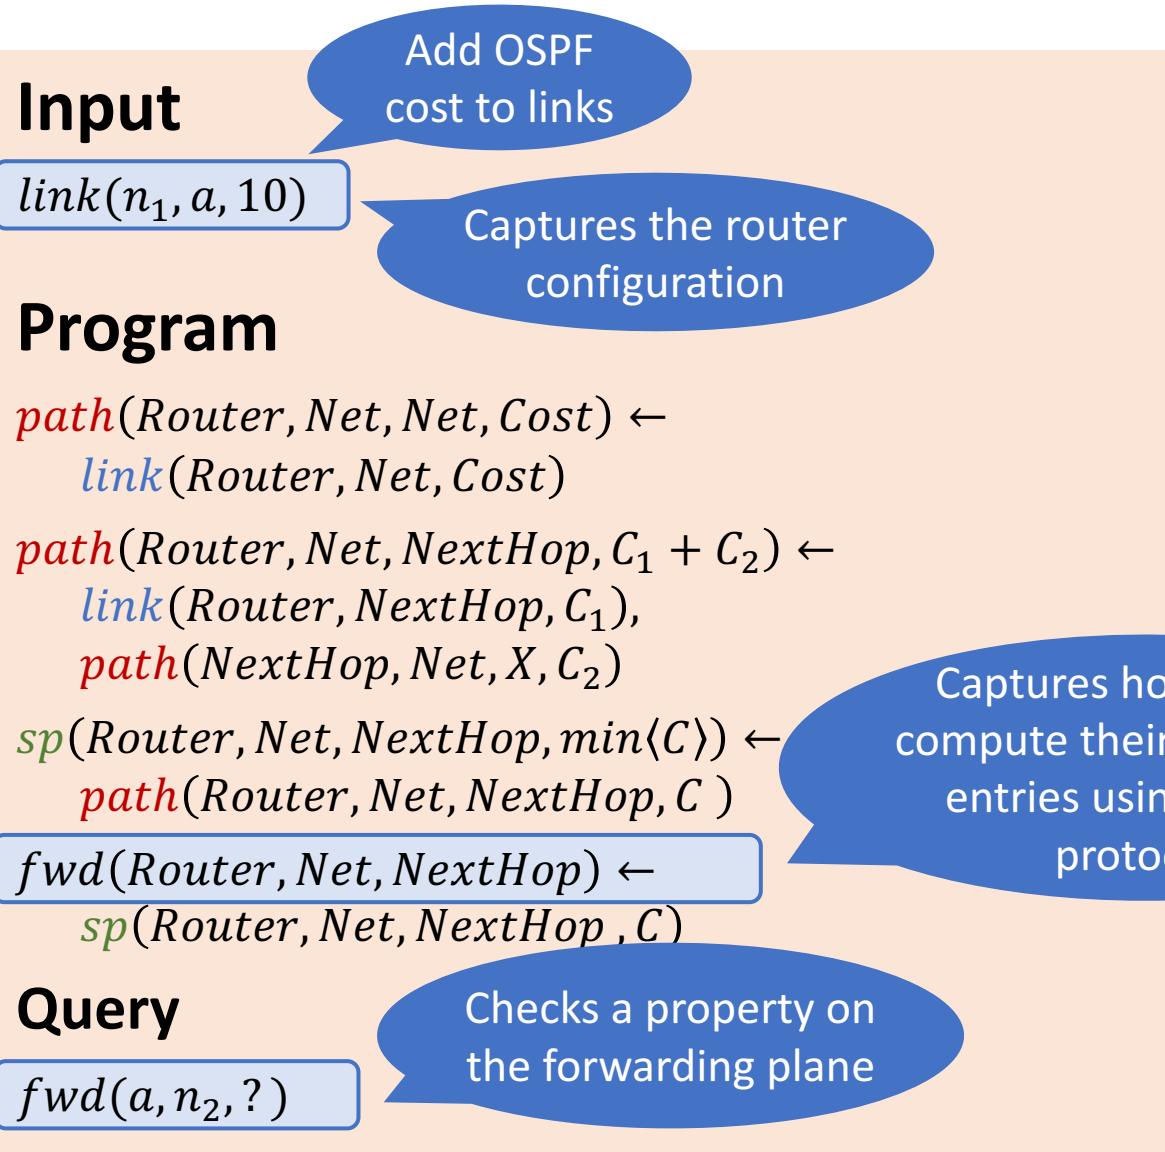
Query

$path(n1, n2)?$



Can we capture the network's forwarding plane?

Datalog (Shortest-path Routing)



Routing Requirements

Paths

Packets for traffic class TC must follow the path

$r_1 \rightarrow \dots \rightarrow r_n$

$$fwd(r_1, tc, r_2) \wedge \dots \wedge fwd(r_{n-1}, tc, r_n)$$

Routing Requirements

Paths

Packets for traffic class TC must follow the path

$$r_1 \rightarrow \dots \rightarrow r_n$$

$$\text{fwd}(r_1, tc, r_2) \wedge \dots \wedge \text{fwd}(r_{n-1}, tc, r_n)$$

Traffic isolation

The paths for two distinct traffic classes tc_1 and tc_2 do not share links in the same direction

$$\forall R_1, R_2. \text{fwd}(R_1, tc_1, R_2) \Rightarrow \neg \text{fwd}(R_1, tc_2, R_2)$$

Routing Requirements

Paths

Packets for traffic class TC must follow the path

$$r_1 \rightarrow \dots \rightarrow r_n$$

$$\text{fwd}(r_1, tc, r_2) \wedge \dots \wedge \text{fwd}(r_{n-1}, tc, r_n)$$

Traffic isolation

The paths for two distinct traffic classes tc_1 and tc_2 do not share links in the same direction

Reachability

Packets for traffic class tc can reach router r_2 from router r_1

$$\text{reach}(r_1, tc, r_2)$$

Routing Requirements

Paths

Packets for traffic class TC must follow the path

$$r_1 \rightarrow \dots \rightarrow r_n$$

$$\text{fwd}(r_1, tc, r_2) \wedge \dots \wedge \text{fwd}(r_{n-1}, tc, r_n)$$

Traffic isolation

The paths for two distinct traffic classes tc_1 and tc_2 do not share links in the same direction

Reachability

Packets for traffic class tc can reach router r_2 from router r_1

$$\text{reach}(r_1, tc, r_2)$$

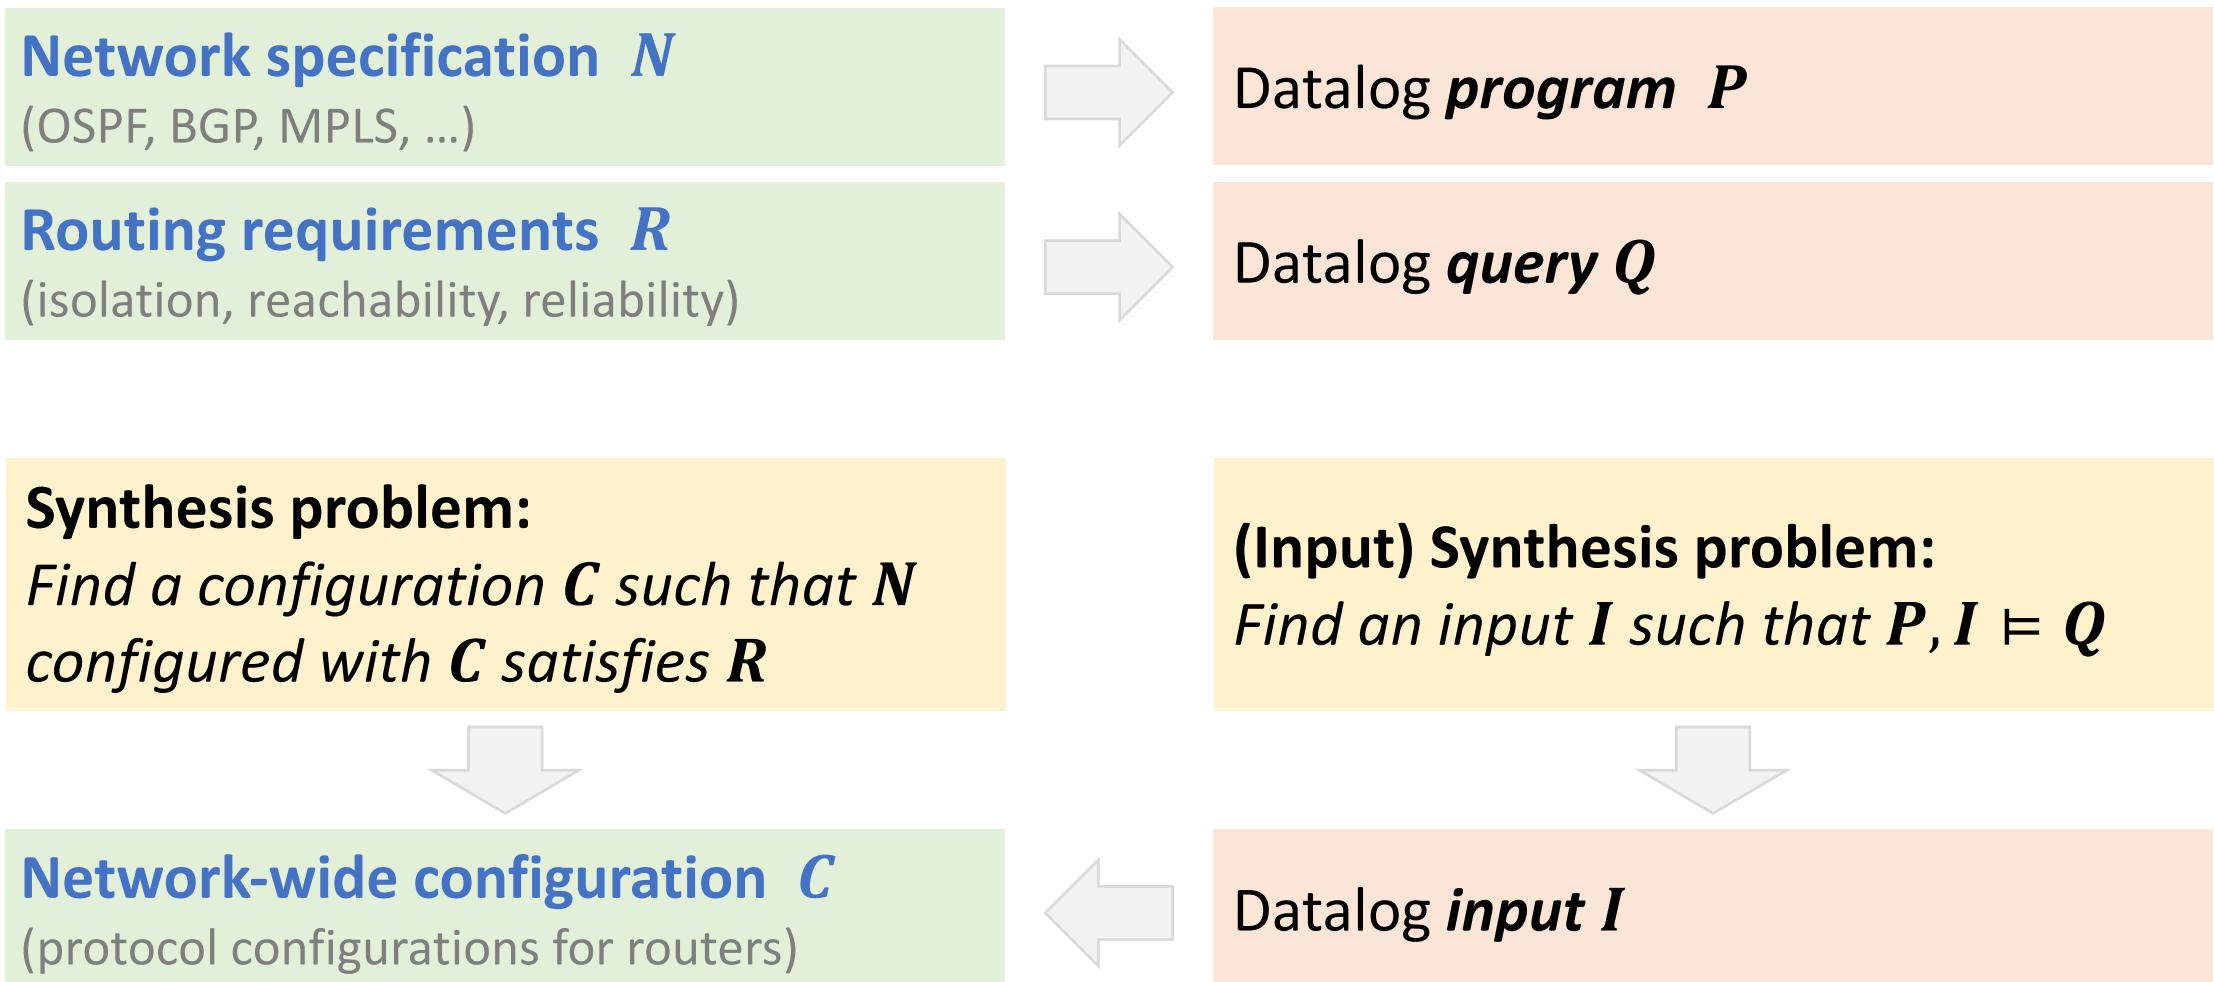
Loop-freeness

The forwarding plane has no loops

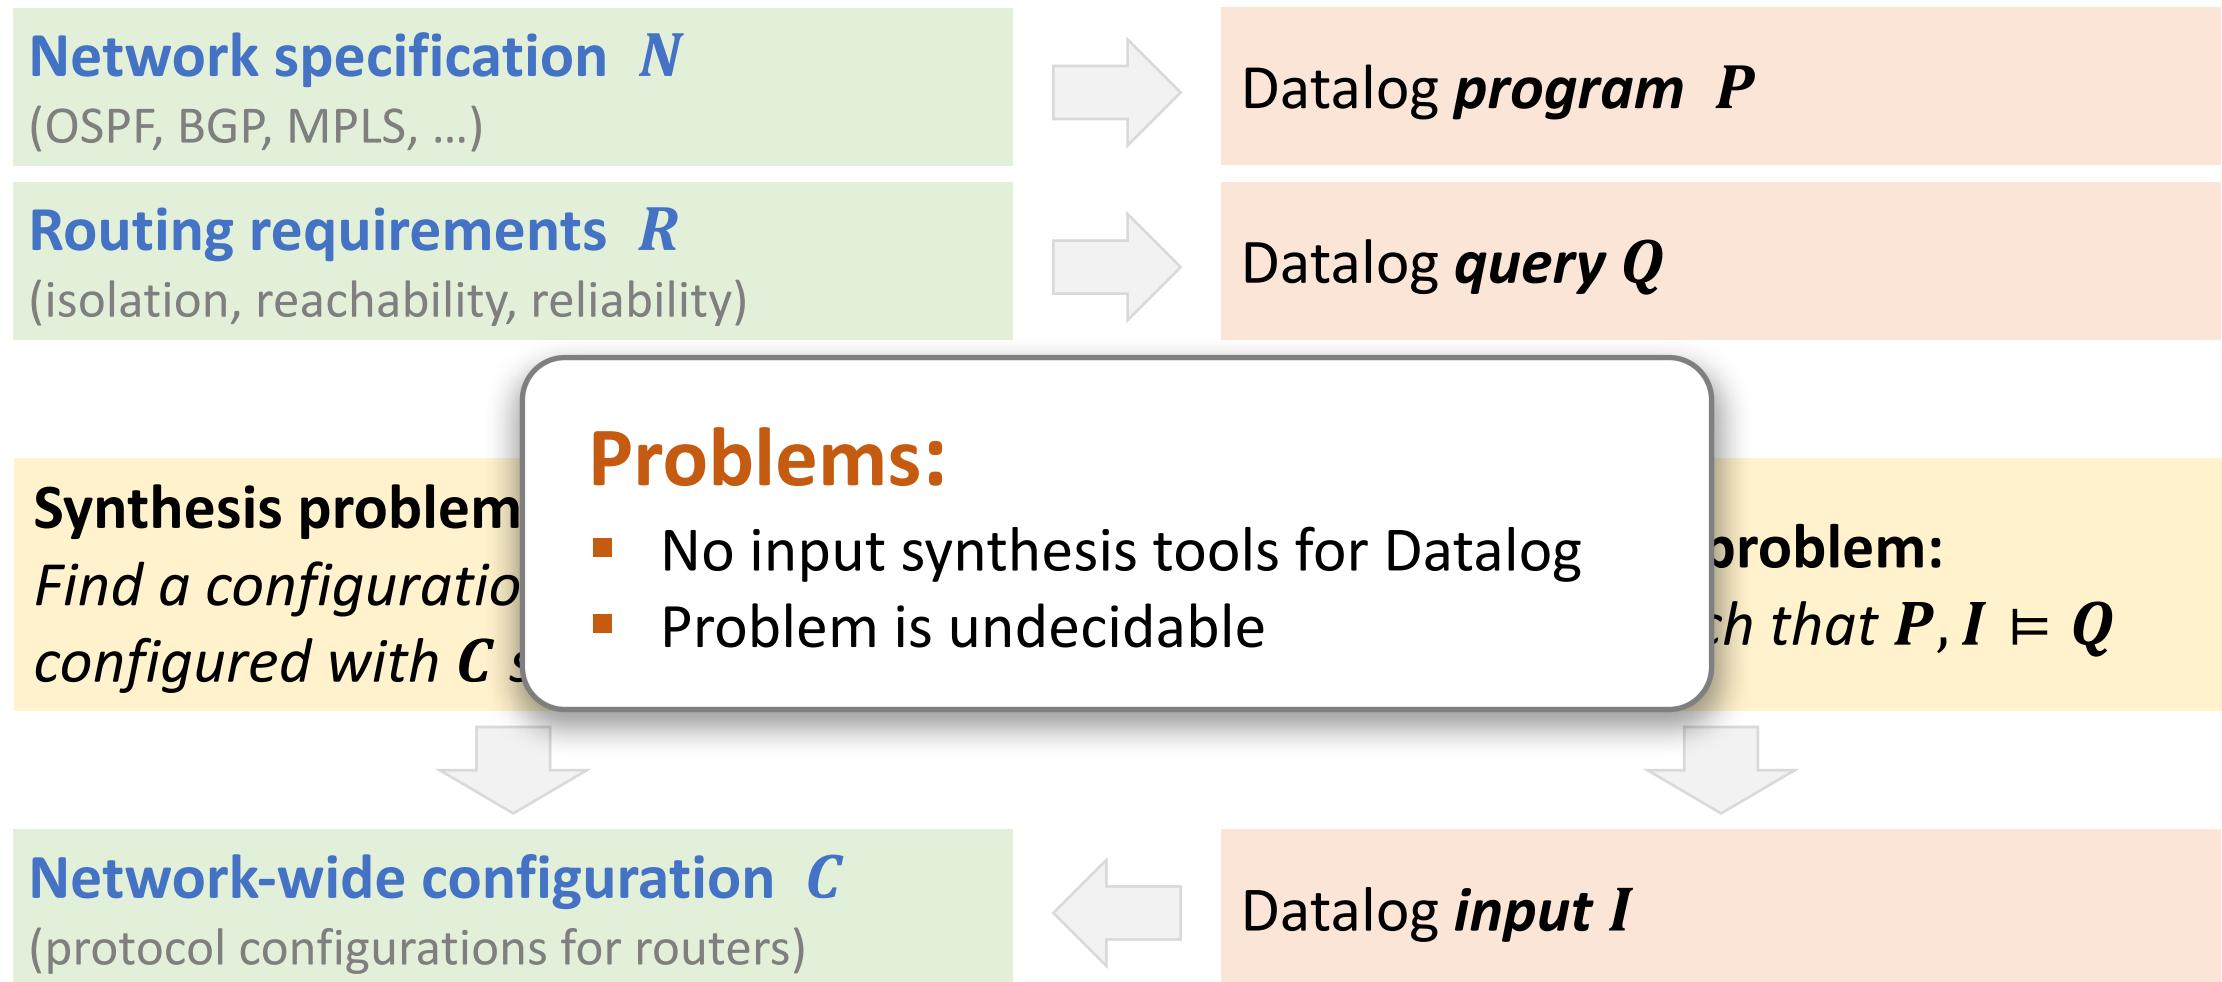
$$\forall TC, R. \neg \text{reach}(R, TC, R)$$

Network-wide Configuration Synthesis

Network-wide Configuration Synthesis



Network-wide Configuration Synthesis



Input Synthesis for Datalog



Key idea: Reduce to solving logical constraints

Input Synthesis for Datalog (via Constraint Solving)

Datalog program P

$$\begin{aligned} path(X, Y) &\leftarrow link(X, Y) \\ path(X, Y) &\leftarrow link(X, Z), path(Z, Y) \end{aligned}$$

Datalog query Q

$$path(a, c) \wedge \neg link(a, c)$$

Generate constraints

$$\forall X, Y. path_1(X, Y) \Leftrightarrow link(X, Y)$$
$$\forall X, Y. path_2(X, Y) \Leftrightarrow \left(link(X, Y) \vee \left(\exists Z. (link(X, Z) \wedge path_1(Z, Y)) \right) \right)$$
$$path_2(a, c) \wedge \neg link(a, c)$$

Bounded unrolling
for recursive rules

Constraints ψ

Constraint solving

$$link(a, b), link(b, c),$$
$$path_1(a, b), path_1(b, c)$$
$$path_2(a, b), path_2(b, c), path_2(a, c)$$

Model $M \models \psi$

Derive input

$$link(a, b), link(b, c)$$

Datalog input I

Input Synthesis for Datalog (via Constraint Solving)

Datalog program P

$$\begin{aligned} path(X, Y) &\leftarrow link(X, Y) \\ path(X, Y) &\leftarrow link(X, Z), path(Z, Y) \end{aligned}$$

Datalog query Q

$$path(a, c) \wedge \neg link(a, c)$$

Generate constraints

$$\forall X, Y. path_1(X, Y) \Leftrightarrow link(X, Y)$$
$$\forall X, Y. path_2(X, Y) \Leftrightarrow \left(link(X, Y) \vee \left(\exists Z. (link(X, Z) \wedge path_1(Z, Y)) \right) \right)$$
$$path_2(a, c) \wedge \neg$$

Bounded unrolling
for recursive rules

Constraints ψ

Theorem: $M \models \psi$ if and only if $P, I \models Q$

$$link(a, b), link(b, c)$$
$$path_1(a, b), path_1(b, c)$$
$$path_2(a, b), path_2(b, c), path_2(a, c)$$

Model $M \models \psi$

Derive input

$$link(a, b), link(b, c)$$

Datalog input I

Generating Constraints

link(N1, A, ?)

...

path(Router, Net, Net, Cost) ← link(Router, Net, Cost)
path(Router, Net, NextHop, C₁ + C₂) ← link(Router, NextHop, C₁), path(NextHop, Net, X, C₂)



$\forall X, Y. \text{path}_1(X, Y) \Leftrightarrow \text{link}(X, Y)$
 $\forall X, Y. \text{path}_2(X, Y) \Leftrightarrow (\text{link}(X, Y) \vee (\exists Z. (\text{link}(X, Z) \wedge \text{path}_1(Z, Y))))$
 $\text{path}_2(a, c) \wedge \neg \text{link}(a, c)$

sp(Router, Net, NextHop, min(C)) ← path(Router, Net, NextHop, C)

fwd(Router, Net, NextHop) ← sp(Router, Net, NextHop , C)

fwd(a, n₂, d) \wedge fwd(a, n₃, b) \wedge ...

Generating Constraints

link(N1, A, ?)

...



Input

path(Router, Net, Net, Cost) ←

link(Router, Net, Cost)

path(Router, Net, NextHop, C₁ + C₂) ←

link(Router, NextHop, C₁),

path(NextHop, Net, X, C₂)



Box 1

sp(Router, Net, NextHop, min{C}) ←

path(Router, Net, NextHop, C)



Box 2

fwd(Router, Net, NextHop) ←

sp(Router, Net, NextHop , C)



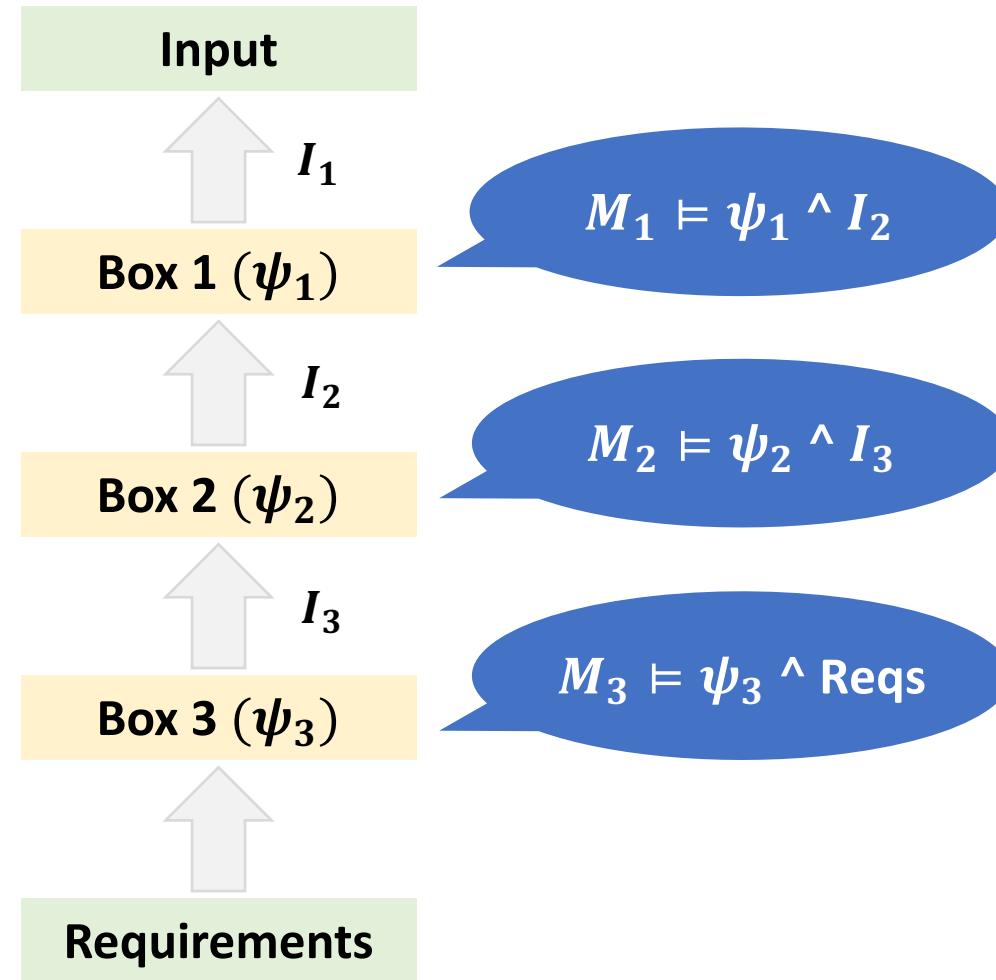
Box 3

fwd(a, n₂, d) ^ fwd(a, n₃, b) ^ ...



Requirements

Synthesis Algorithm (Naïve Approach)



Synthesis Algorithm (Naïve Approach)

Naïve approach does NOT work

1. Not all valid SMT models are valid network configurations.
2. The search space of valid inputs is huge.

Synthesis Algorithm (Our Approach)

Input Constraints

Ensure the synthesized configurations are valid

Ex. OSPF Link costs are positive numbers

Synthesis Algorithm (Our Approach)

Input Constraints

Ensure the synthesized configurations are valid

Ex. OSPF Link costs are positive numbers

Domain specific constraints

Reduce the search space

Ex. Routers only forward to neighbours

Synthesis Algorithm (Our Approach)

Input Constraints

Ensure the synthesized configurations are valid

Ex. OSPF Link costs are positive numbers

Domain specific constraints

Reduce the search space

Ex. Routers only forward to neighbours

Learn from mistakes

Don't reuse the same invalid input

$$M_{boxi} \models \psi \wedge l_{i+1} \wedge \neg l_i$$

Synthesis Algorithm (Our Approach)

Input Constraints

Ensure the synthesized configurations are valid

Ex. OSPF Link costs are positive numbers

Domain specific constraints

Reduce the search space

Ex. Routers only forward to neighbours

Learn from mistakes

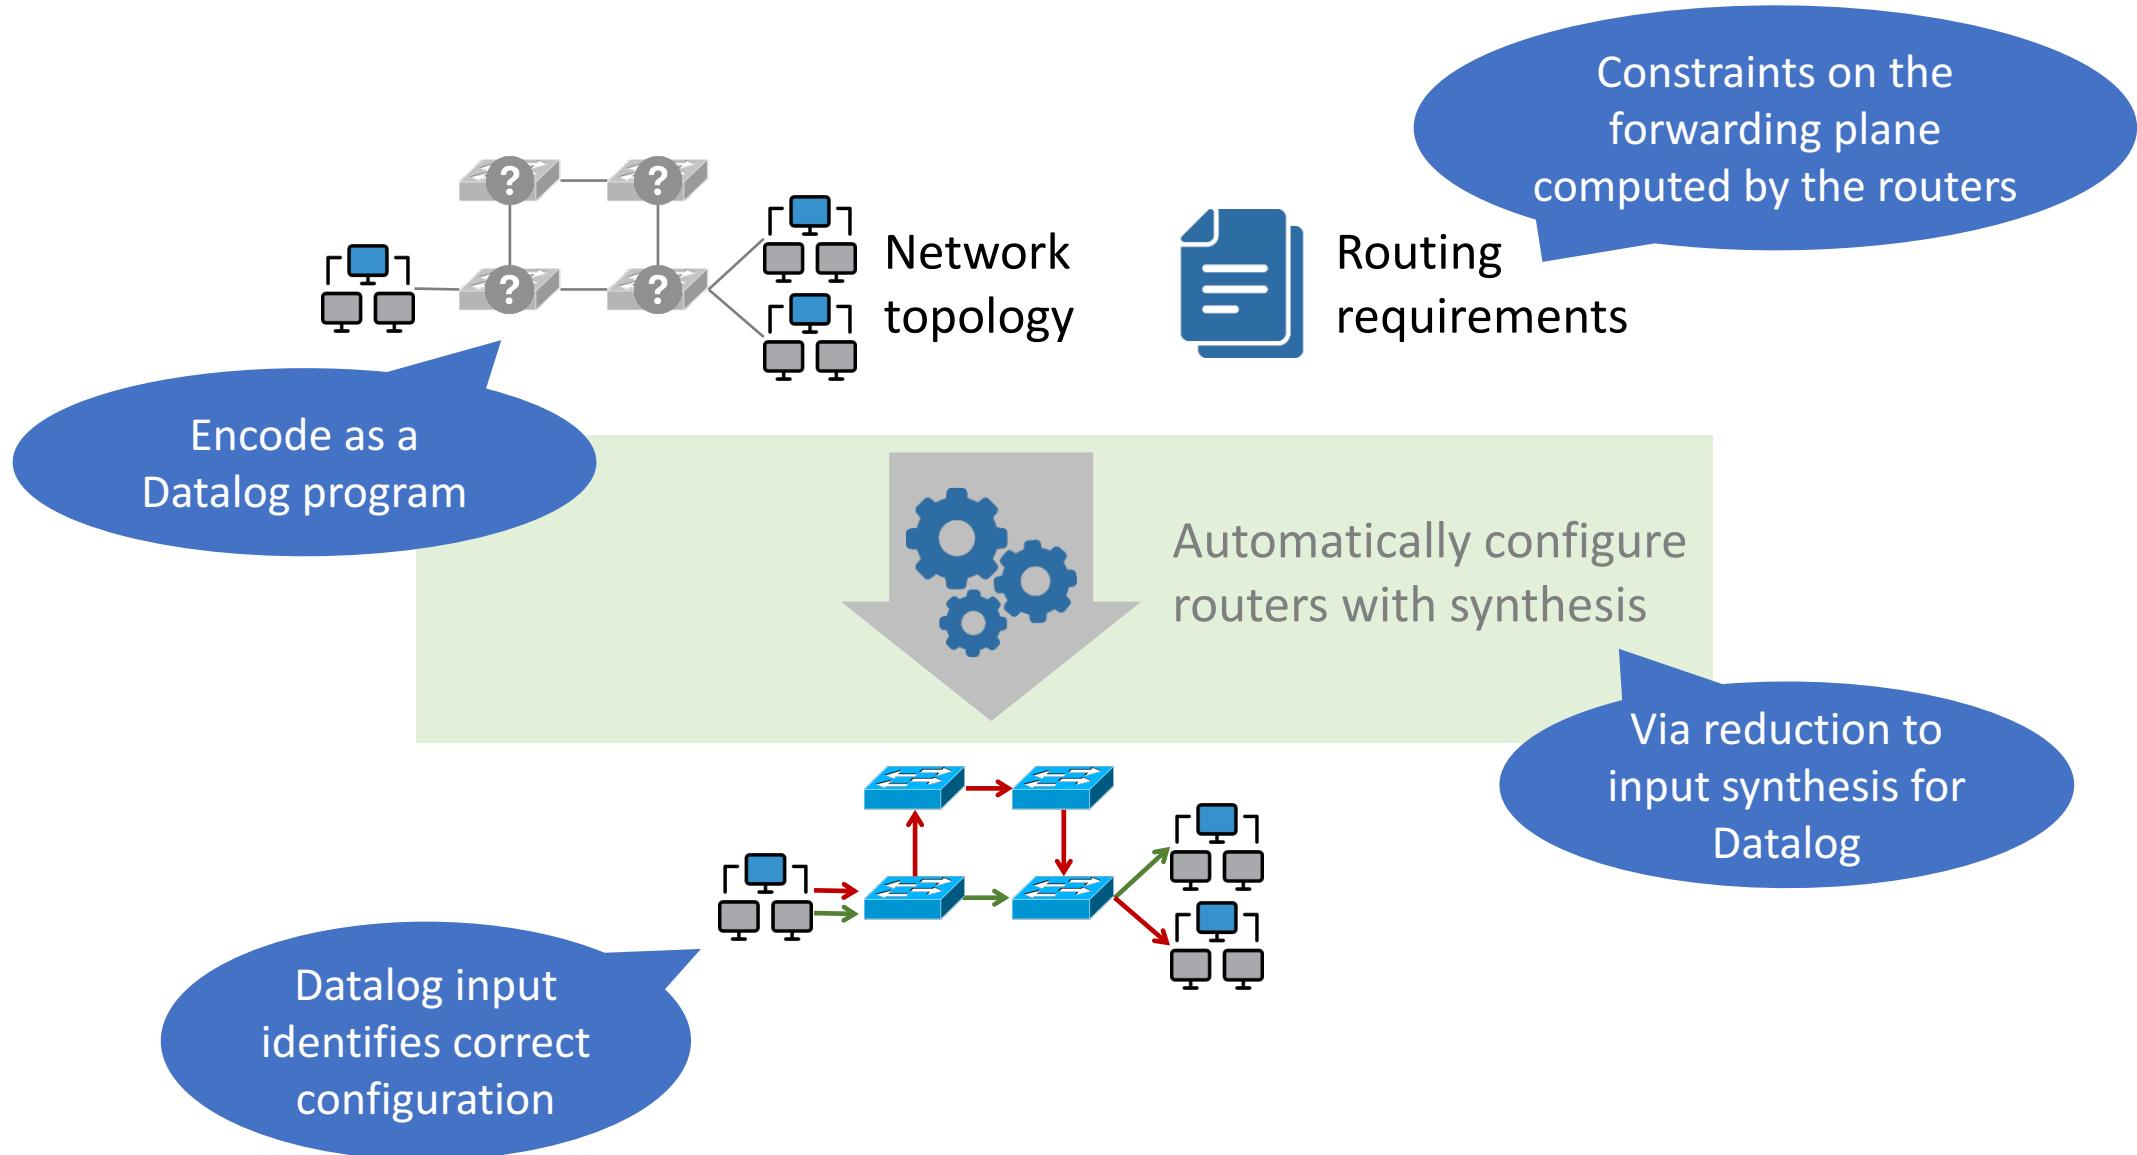
Don't reuse the same invalid input

$$M_{boxi} \models \psi \wedge l_{i+1} \wedge \neg l_i$$

Partial evaluation

Don't compute already known values

Network-wide Configuration Synthesis



Implementation

Implementation

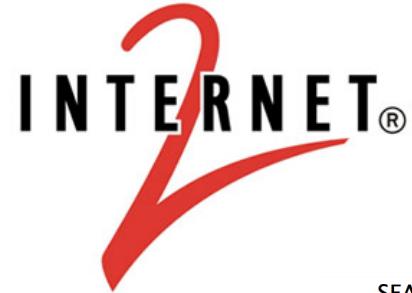
The SyNET system (<http://synt.ethz.ch>)

- Written in **Python** ($\approx 4K$ lines of code)
- Protocols encoded in **stratified Datalog** (≈ 100 rules)
- Uses the **Z3** constraint solver
- Outputs **CISCO** configurations
- Supports **BGP**, **OSPF**, and **static routes**

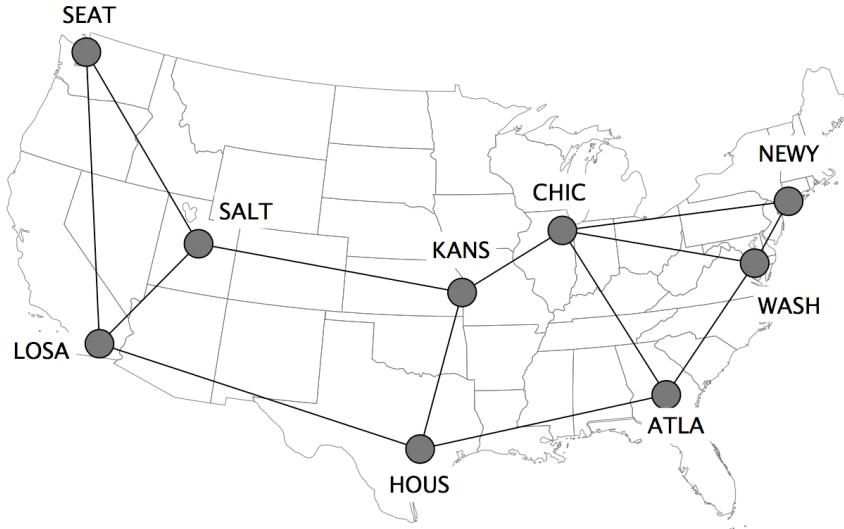
```
! A snippet from router A
interface f0/1
    ip address 10.0.0.2 255.255.255.254
    ip ospf cost 10
    description "To B"
interface f0/0
    ip address 10.0.0.0 255.255.255.254
    ip ospf cost 65530
    description "To C"
interface f1/0
    ip address 10.0.0.4 255.255.255.254
    ip ospf cost 65530
    description "To D"
!
```

CISCO configuration
output by SyNET

Experiments



Experiment



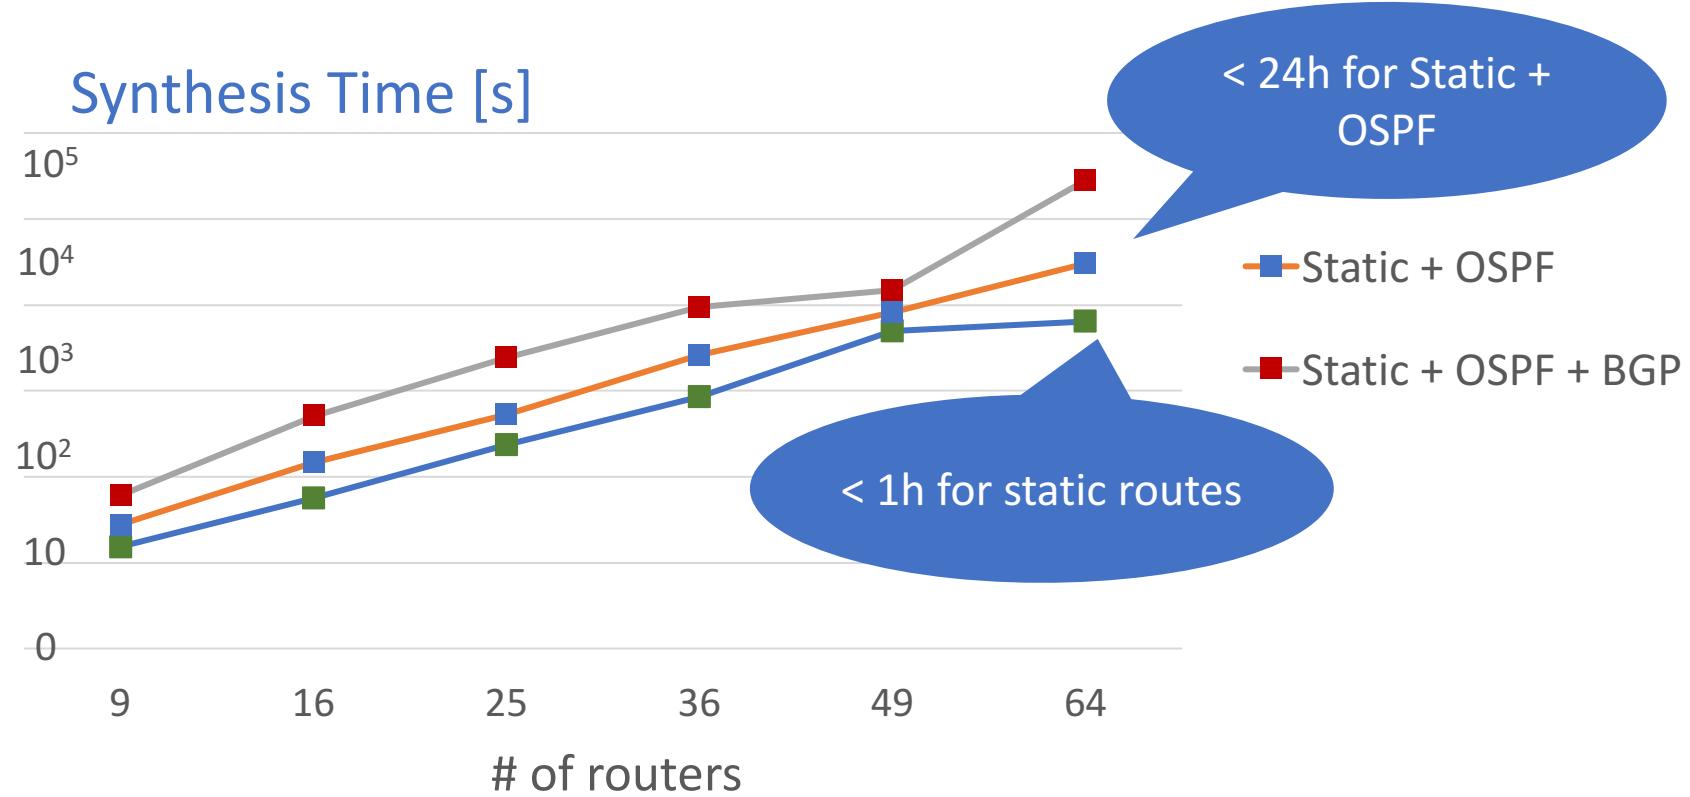
US-based network connecting major universities and research institutes

Protocols / # Traffic classes	1 class	5 classes	10 classes
Static	1.3s	2.0s	4.0s
Static + OSPF	9.0s	21.3s	49.3s
Static + OSPF + BGP	13.3s	22.7s	1m19.7s

Synthesis Times

Scalability Experiment

- Grid topologies with up to 64 routers
- Requirements for 10 traffic classes



Summary: Programmable Networks with Synthesis

